

Conference on Health, Inference, and Learning (CHIL) 2025

Xuhai Orson Xu*

Columbia University, New York, NY, USA

XX2489@COLUMBIA.EDU

Edward Choi*

Korea Advanced Institute of Science & Technology (KAIST), Daejeon, South Korea

EDWARDCHOI@KAIST.AC.KR

Pankhuri Singhal*

University of Pennsylvania, Philadelphia, PA, USA

SINGHALP@PENNMEDICINE.UPENN.EDU

Walter Gerych*

Massachusetts Institute of Technology, Cambridge, MA, USA

WGERYCH@MIT.EDU

Shengpu Tang*

Emory University, Atlanta, GA, USA

SHENGPU.TANG@EMORY.EDU

Monica Agrawal*

Duke University, Durham, NC, USA

MONICA.AGRAWAL@DUKE.EDU

Adarsh Subbaswamy*

US Food and Drug Administration, Silver Spring, MD, USA

ADARSH.SUBBASWAMY@FDA.HHS.GOV

Elena Sizikova*

US Food and Drug Administration, Silver Spring, MD, USA

ELENA.SIZIKOVA@FDA.HHS.GOV

Jessilyn Dunn*

Duke University, Durham, NC, USA

JESSILYN.DUNN@DUKE.EDU

Roxana Daneshjou*

Stanford University, Stanford, CA, USA

ROXANAD@STANFORD.EDU

Tasmie Sarker*

Association for Health Learning and Inference, New York, NY, USA

TASMIE.SARKER@AHLI.CC

Irene Y. Chen*

University of California, Berkeley, Berkeley, CA, USA

IYCHEN@BERKELEY.EDU

Matthew McDermott*

Harvard Medical School, Boston, MA, USA

MATTMCDERMOTT8@GMAIL.COM

Abstract

The Conference on Health, Inference, and Learning (CHIL), brings together a cross-disciplinary group of clinicians and researchers, from industry and academia, with the goal of advancing machine learning for health. CHIL has been an official conference of the Association for Health Learning and Inference (AHLI) since 2022. This volume contains proceedings of the sixth annual CHIL conference, held at the University of California, Berkeley in the US.

1. Introduction

The Conference on Health, Inference, and Learning (CHIL) brings together a cross-disciplinary community of clinicians and researchers from both academia and industry, with a shared mission to advance machine learning for health ([Conference on Health, Inference, and Learning \(CHIL\), 2025](#)). CHIL serves as a platform for this community to connect, exchange ideas, and explore emerging directions in the field.

In recent years, machine learning and “AI”—especially generative models—have moved into the spotlight, capturing mainstream attention and opening up new possibilities for health applications. As excitement grows, it is more important

* These authors contributed equally

than ever for researchers and practitioners to stay grounded: to pursue innovation responsibly while maintaining a strong focus on methods and systems that are safe, fair, and ultimately beneficial to patients and care providers. Meeting this challenge demands collaboration across diverse domains, including clinical practice, computer science, and the social sciences.

Through CHIL, we aim to foster these vital collaborations and accelerate meaningful progress in machine learning for health. This volume includes the proceedings of the sixth annual CHIL conference, held at the University of California, Berkeley.

1.1. Organizing Committee

Matthew McDermott (General Chair), Irene Chen (General Chair), Jessilyn Dunn (Program Chair), Roxana Daneshjou (Program Chair), Vicky Bikia (Unconference Chair), Kaveri Thakoor (Research Roundtable Chair), Hayoung Jeong (Research Roundtable Chair), Emma Rocheteau (Doctoral Symposium Chair), Matt Engelhard (Doctoral Symposium Chair), Edward Choi (Proceedings Chair), Pankhuri Singhal (Proceedings Chair), Xuhai “Orson” Xu (Proceedings Chair), Walter Gerych (Senior Area Chair, Track 1), Shengpu Tang (Senior Area Chair, Track 2), Adarsh Subbaswamy (Senior Area Chair, Track 2), Monica Agrawal (Senior Area Chair, Track 3), Elena Sizikova (Senior Area Chair, Track 3), Xinyang Han (Logistics Chair, Local), Radhika Bhalerao (Logistics Chair, Local), Kai Wang (Finance Chair), Ahmed Alaa (Finance Chair).

2. Program

CHIL’s program was designed to showcase exceptional work through invited talks, to open opportunities for debate on topical issues, and to promote collaboration and learning through roundtables, lightning talks, posters, and a doctoral symposium.

2.1. Keynote Speakers

Our keynote speakers were selected to reflect a broad range of perspectives and research directions in machine learning for health:

Maia Hightower: Maia Hightower is the CEO and Founder of Equality AI. Her work focuses on responsible AI, digital health equity, and healthcare innovation.

Ida Sim: Ida Sim is a Professor of Medicine at the University of California, San Francisco. Her research lies at the intersection of digital health, clinical informatics, and computational precision health.

Guillermo Sapiro: Guillermo Sapiro is the Augustine Family Professor of Electrical and Computer Engineering at Princeton University and a Distinguished Engineer with Apple. His research spans machine learning, computer vision, and biomedical imaging.

Judy Gichoya: Judy Gichoya is an Associate Professor of Radiology and Imaging Sciences at Emory University. Her work focuses on artificial intelligence in medical imaging, health equity, and informatics.

Peter Szolovits: Peter Szolovits is a Professor of Computer Science and Engineering at the Massachusetts Institute of Technology. He pioneers research in artificial intelligence, clinical decision-making, and biomedical informatics.

2.2. Unconference

The CHIL unconference is a dynamic, participant-driven forum designed to foster discussion, collaboration, and knowledge sharing on pressing issues in machine learning for health. The 2025 unconference featured opening remarks, a panel on AI entrepreneurship in health, a series of lightning talks, two topical panel discussions, and multiple parallel roundtable sessions. The day concluded with closing reflections and a networking reception.

2.3. Panels

Our panel sessions explore topical and evolving issues in machine learning for health, featuring leading voices across academia, healthcare, and public research institutions:

GenAI Use and Ethics in Health: Jenna Lester (UCSF), Leo Celi (MIT), Pranav Rajpurkar (Harvard Medical School), Deborah Raji (UC Berkeley). *Moderated by Emily Alsentzer (Stanford University).*

Translational Research in a Rapidly Changing World: Josh Fessel (NIH, former), Nigam Shah (Stanford University), Chris Lunt (NIH), Jean Feng (UCSF). *Moderated by Peter Washington (UCSF).*

2.4. Research Roundtables

Our roundtables offer interactive discussions around key challenges and future directions in ML for health.

Each session is co-led by domain experts and encourages participation from attendees across disciplines: (1) Explainability, Interpretability, and Transparency — Ashery Christopher Mbilinyi, Chanwoo Kim; (2) Explainability/Interpretability in Sequence Models — Tanveer Syeda-Mahmood, Shashank Yadav; (3) Uncertainty, Bias, and Fairness — Emma Pierson, Aparajita Kashyap; (4) Causality — Marie-Laure Charpignon, Nitish Nagesh; (5) Scalable AI Solutions: Domain Adaptation — Aishwarya Mandyam, Simon Lee; (6) Scalable AI Solutions: Foundation Models — Akane Sano, Arvind Pillai; (7) Scalable AI Solutions: Learning from Small Medical Data — Emily Alsentzer, Brighton Nuwagira; (8) Multimodal Methods — Jason Fries, Bill Chen; (9) Scalable and Translational Healthcare Solutions — Niharika S. DSouza, Yixing Jiang.

2.5. Fireside Chat

Andrea Downing is a leading patient advocate and co-founder of The Light Collective, a nonprofit focused on data rights and privacy for patient communities. Together, Roxana and Andrea will shape a conversation centered on the patient perspective in healthcare AI, highlighting challenges, opportunities, and ethical considerations in patient-centered innovation.

2.6. Year In Review

The Year in Review session is a curated overview of influential research published in the past year that aligns with the CHIL community’s mission. Its goal is to highlight emerging themes, impactful contributions, and noteworthy trends across clinical machine learning and health informatics. By surfacing high-quality papers from a wide range of venues, the session aims to foster discussion, reflection, and inspiration for future work within the CHIL community.

3. Papers: Guidelines and Selection

Papers for the proceedings were solicited across three tracks at the intersection of health and machine learning: *Models and Methods*; *Applications and Practice*; and *Impact and Society*. The submission and review process was entirely managed through the OpenReview platform ([OpenReview, 2025](#)).

3.1. Submissions

We received a total of 104 submissions across the three tracks. All submissions were reviewed by a minimum of four reviewers, except where the submission was subsequently desk rejected ($n=3$) or withdrawn ($n=5$).

3.2. Peer review

CHIL follows a mutually anonymous (“double blind”) review process with five broad phases:

1. Reviewer bidding
2. Assignment of papers to reviewers
3. Completion of reviews
4. Discussion, and author rebuttals
5. Decision and notification

Reviewers were asked to consider aspects of relevance, quality, originality, clarity, and significance when preparing their reviews. After reviews were completed, a discussion phase allowed open communication between authors, reviewers, and chairs. Area Chairs then created meta-reviews and recommendations based on the reviews. Final decisions were made by Senior Area Chairs in discussion with Area Chairs, Proceedings Chairs, and reviewers, where appropriate.

4. Proceedings

Forty-two submissions (40%) were accepted for publication in the archived conference proceedings. Some prominent themes amongst the accepted papers include:

- Generative AI and Foundation Models: [Im et al.](#), [Kyung et al.](#), [Lin et al.](#), [Pillai et al.](#)
- Causal Inference for Clinical Insights: [Chen et al.](#), [Mukherjee et al.](#), [Ji et al.](#), [Noroozizadeh et al.](#), [Sunog et al.](#), [Jiang et al.](#)
- Reliability, Fairness, and Interpretability: [Chiang et al.](#), [Kapash et al.](#), [Yadav and Subbian](#), [Jin et al.](#), [Jo et al.](#), [Liang et al.](#), [Yun et al.](#), [Oh and Bui](#), [Yan et al.](#), [Jin and Ghassemi](#), [Lopez et al.](#), [Yu et al.](#), [Jia et al.](#)
- Multimodal Learning and Data Integration: [El-hussein et al.](#), [Tang et al.](#), [Wang et al. \(a\)](#), [Yang et al.](#), [Liu et al.](#)

- Time Series, Wearables, and Biosignals: Wang et al. (b), Potosnak et al., Stemmer et al., Kiafar et al., Moghaddami et al.
- Benchmarking, Evaluation, and Methods for Clinical Translation: Zhang et al. (b), Park et al., Kraft et al., Toye et al., Zhang et al. (a), Acosta et al., Turura et al., Fong and Motani, Elsharief et al.

4.1. Awards

Three papers were selected for a Best Paper Award across the conference tracks.

4.1.1. MODELS AND METHODS

Multi-View Contrastive Learning for Robust Domain Adaptation in Medical Time Series Analysis by YongKyung Oh, Alex Bui (Oh and Bui)

4.1.2. APPLICATIONS AND PRACTICE

LabTOP: A Unified Model for Lab Test Outcome Prediction on Electronic Health Records by Sujeong Im, Jungwoo Oh, Edward Choi (Im et al.).

4.1.3. IMPACT AND SOCIETY

Learning Disease Progression Models That Capture Health Disparities by Erica Chiang, Divya M Shanmugam, Ashley Beecy, Gabriel Sayer, Deborah Estrin, Nikhil Garg, Emma Pierson (Chiang et al.).

5. CHIL 2026 in Seattle, WA, USA

CHIL 2026 will be held in Seattle, WA, USA in June. Save the date! The conference will be led by general chairs Irene Chen and Emily Alsentzer. If you would like to volunteer at the conference or serve on our committee, please contact us at info@chilconference.org.

Acknowledgments

We thank the Association for Health Learning and Inference (AHLI) for their continued support.

5.1. Sponsors

We are deeply grateful for the generous support of our sponsors, whose contributions made CHIL 2025 possible and continue to propel the community at the intersection of health, inference, and learning.

Diamond Sponsor: we extend our special thanks to our Diamond-level sponsor, the Computational Precision Health Program at UCSF / UC Berkeley. Their support and leadership have been foundational to the success of this year’s conference.

Gold Sponsors: we are also proud to recognize our Gold-level sponsors—ADIA Lab, Google, Schmidt Sciences, and Optum—for their strong commitment to advancing research and innovation in health and machine learning.

Silver Sponsors: our Silver-level sponsors—GSK plc, Apple, and Silicon Valley Bank—provided valuable support that helped bring this event to life.

Startup Sponsors: we also warmly thank our Startup-level sponsors, Stanford Health Care and Define Ventures, for their enthusiastic engagement from academia and venture capital.

The collective generosity and support of all our sponsors has been instrumental in making CHIL 2025 a success.

5.2. Reviewers

The proceedings would not have been possible without the diligent and generous work of our reviewers.

5.2.1. LIST OF SENIOR AREA CHAIRS

Walter Gerych, Shengpu Tang, Adarsh Subbaswamy, Monica Agrawal, and Elena Sizikova.

Dr. Subbaswamy and Dr. Sizikova stepped down and transferred their roles to Dr. Tang and Dr. Agrawal during the paper submission stage due to administrative constraints, ensuring a smooth continuation of the committee’s responsibilities.

5.2.2. LIST OF AREA CHAIRS

Jean Feng, Zehang Li, Michael Oberst, Intae Moon, Jessica Dafflon, Sarah Tan, Megan Coffee, Stephen R Pfohl, Julia E Vogt, Samantha Kleinberg, Zepeng Frazier Huo, Ayah Zirikly, Kimia Ghobadi, and Weimin Zhou.

5.2.3. LIST OF REVIEWERS

Yue Guo, Kamilė Stankevičiūtė, Tianhao Li, Bharath Kumar Varma Sagi, Chunyi Wang, Kristen A Severson, Sandeep Angara, Andre Stollenwerk, Haochen Sui, Farzana Islam Adiba, William Han, Xian Wu, Chaoyi Tan, Taylor W. Killian, Mingxuan Liu, Vasanth Rajendran, Chuizheng Meng, Zefang Liu, Xinyu Xiong, Evgenii Evstafev, Anna Zink, Junu Kim, Botao Zhang, Daeun Kyung, Changxin Xu, Sicong Huang, Cheng Fei, Kaiyu Mu, Arkin Gupta, Xinyang Han, Mahak Shah, Fabian Gröger, Mehak Gupta, Aparajita Saraf, Huixin Zhan, Eptehal Nashnoush, Ketong Shao, Zhun Zhou, Jungwoo Oh, Zhengyu Fang, Harsh Sharma, Michael S Yao, Chandrasekar Ramachandran, Lovedeep Gondara, Mehul Motani, Preetish Rath, Cédric Wemmert, Yinzhu Quan, Alex Fedorov, Jichan Chung, Mingyang Wan, Naman Goyal, Rakib Islam, Srihari Jayakumar, Marta Avalos, Dhruv Naik, Ariel Guerra-Adames, Trenton Chang, Ashley Wu, Srijan Mishra, Aishwarya Jadhav, Ajitesh Jain, Harlin Lee, Sheshananda Reddy Kandula, Ziyue Yu, Shahriar Noroozizadeh, Harsh Nilesh Pathak, Chiyu Cheng, Samarth Agrawal, Esma Yildirim, Naveen Janaki Raman, Yuan Tian, Mollie Marian Mckillop, Marie-Laure Charpignon, Shaeke Salman, Stefan Hegselmann, Charles B. Delahunt, Gaurang Deshpande, Wangzhi Dai, Yingxin Lai, Luckyson Khaidem, Jiho Kim, Hariharan Ragothaman, Divyam Madaan, Ayush Noori, Yuwen Chen, Brighton Nuwagira, Yao Su, Gene-Ping Yang, Jiahe Xu, Fengrui Zhang, Quanchao Lu, Xiaobin Shen, Zhuolin Hao, Zhe Huang, Ashery Mbilinyi, Keith Trnka, Akaash Vishal Hazarika, Benedikt Schesch, Shengqi Jian, Jiaqing Zhang, Vasiliki Bikia, Qisen Cheng, Siwei Wang, Mark Braunstein, Yining Liu, Rui Yang, Vasundhara Agrawal, Anish Mahishi, Yue Li, Jiajing Chen, Hongzhou Luan, Vin Bhaskara, Chen Chen, Yinghao Zhu, Luna Zhang, David Chen, Yong-Yeon Jo, Maryam Khalid, Xiaopo Cheng, Zijian Zhang, Melissa Danielle McCradden, Benjamin C Warner, Hyungyung Lee, Wian Stipp, Stephanie Hyland, Zhen Huang, Kaitian Jin, Wanying Dou, Vibhu Verma, Bobak J Mortazavi, Sumit Mukherjee, Xu Cao, Naoki Nonaka, Nilesh Pandey, Junyi Gao, Lorenzo A. Rossi, Madiha Shakil Mirza, Shuo Li, Zongliang Ji, Feng Wu, Prateek Chhikara, Salvatore Tedesco, Pallavi Gudipati, Kyunghoon Hur, Yifei Chen, Abdullah Mamun, Xiaozhou Shi, Yuhan Li, Wayne Isaac Tan Uy, Emma Charlotte Rocheteau, Xiao Fan, Elizabeth Healey, Lucía Prieto Santamaría,

Sunny Dhamnani, Shreya Gupta, Aishi Huang, Jue Xiao, Bhaskar Goyal, Li Zhang, Jinghan Cao, Botao Zhang, Keshu Wu, Prudhvi Nethi, Chi Xie, Ghulam Ahmed Ansari, Madhulekha Arunmozhi, Bala Siva Sai Akhil Malepati, Ethan Harvey, Jialu Wang, Qiyuan Tian, Yike Zhang, Aryan Jadon, Arnav Saxena, Pranav Kulkarni, Venkat Nilesh Dadi, Mirza Farhan Bin Tarek, Yingrui Ji, Jennifer Yu, Raju Dandigam, Vishal Shah, Rongqi Fan, Tong Xia, Yue Yin, Yuan Zhao, Dayu Hu, Chloe Zhu, Peijie Qiu, Simon A. Lee, Mukesh Yadav, Chang Zhou, Seongjun Yang, Anurag Jayant Vaidya, Zhiyuan Wang, Izzy Chaiken, Shalini Saini, Haoze Wu, Muhammad Arbab Arshad, Karanbir Singh, Frank Rudzicz, Mohammad Amir Salari, Rohit Vashisht, Ali Septiandri, Aishwarya Mandyam, Jiayi Yang, Matthew M. Engelhard, Rahmatollah Beheshti, Songlin Hou, Yungang Li, Megan Coffee, Hyewon Jeong, Ashly Joseph, Shayan Fazeli, Jieshi Chen, Ajantha Ramineni, Sirui Ding, Abhai Pratap Singh, Lili Wu, Mudit Mangal, Bret Nestor, Srinivasarao Daruna, Ismael Villanueva-Miranda, Rajesh Vayyala, Cassandra Parent, Haoran Zhang, and Aparna Balagopalan.

References

- Julian Nicolas Acosta, Xiaoman Zhang, Siddhant Dogra, Hong-Yu Zhou, Seyedmehdi Payabvash, Guido J. Falcone, Eric Karl Oermann, and Pranav Rajpurkar. Headct-one: Enabling granular and controllable automated evaluation of head ct radiology report generation. pages 649–671.
- Winston Chen, Trenton Chang, and Jenna Wiens. Conditional front-door adjustment for heterogeneous treatment assignment effect estimation under non-compliance. pages 194–230.
- Erica Chiang, Divya M Shanmugam, Ashley Beecy, Gabriel Sayer, Deborah Estrin, Nikhil Garg, and Emma Pierson. Learning disease progression models that capture health disparities. pages 1–29.
- Conference on Health, Inference, and Learning (CHIL). CHIL Conference, 2025. URL <https://chilconference.org/>. Accessed: 2025-06-18.
- Ahmed Elhoussein, Paul Meddeb, Abigail Newbury, Jeanne Mirone, Martin Stoll, and Gamze Gursöy. Keep: Integrating medical ontologies with clinical data for robust code embeddings. pages 43–62.

- Shaza Elsharief, Saeed Shurrah, Baraa Al Jorf, Leopoldo Julian Lechuga Lopez, Krzysztof J. Geras, and Farah E. Shamout. Medmod: Multi-modal benchmark for medical prediction tasks with electronic health records and chest x-ray scans. pages 781–803.
- Kei Sen Fong and Mehul Motani. Multi-objective fine-tuning of clinical scoring tables: Adapting to variations in demography and data. pages 744–780.
- Sujeong Im, Jungwoo Oh, and Edward Choi. Labtop: A unified model for lab test outcome prediction on electronic health records. pages 817–843.
- Zongliang Ji, Andre Carlos Kajdacsy-Balla Amaral, Anna Goldenberg, and Rahul G Krishnan. Exosito: Explainable off-policy learning with side information for intensive care unit blood test orders. pages 337–368.
- Furong Jia, David Sontag, and Monica Agrawal. How does my language model understand clinical text? pages 720–743.
- Muyan Jiang, Yunkai Zhang, and Anil Aswani. Learning interactions between continuous treatments and covariates with a semiparametric model. pages 691–707.
- Qixuan Jin and Marzyeh Ghassemi. A study of artifacts on melanoma classification under diffusion-based perturbations. pages 844–861.
- Yeping Jin, Lauren Wise, and Ioannis Paschalidis. Distributionally robust learning in survival analysis. pages 369–380.
- Yong-Yeon Jo, Byeong Tak Lee, Jeong-Ho Hong, Hak Seung Lee, Joon-myung Kwon, and Beom Joon Kim. Test-time calibration: A framework for personalized test-time adaptation in real-world biosignals. pages 381–394.
- Daniel Kapash, Noam Barda, Omer Reingold, Noa Dagan, and Ran Balicer. Multiaccuracy for subpopulation calibration over distribution shift in medical prediction models. pages 130–144.
- Behdokht Kiafar, Pinar Kullu, Rakshith Lokesh, Amit Chaudhari, Qile Wang, Shayla Sharmin, Sagar M. Doshi, Elham Bakhshipour, Erik Thostenson, Joshua Cashaback, and Roghayeh Leila Barmaki. Feasibility of immersive virtual reality and customized robotics with wearable sensors for upper extremity training. pages 543–556.
- Stefan Kraft, Andreas Theissler, Dr. Vera Wienhausen-Wilke, Philipp Walter, Gjergji Kasneci, and Hendrik Lensch. Alpec: A comprehensive evaluation framework and dataset for machine learning-based arousal detection in clinical practice. pages 395–429.
- Daeun Kyung, Junu Kim, Tackeun Kim, and Edward Choi. Towards predicting temporal changes in a patient’s chest x-ray images based on electronic health records. pages 247–267.
- Zhongyuan Liang, Arvind Suresh, and Irene Y. Chen. Treatment non-adherence bias in clinical machine learning: A real-world study on hypertension medication. pages 430–442.
- Yihan Lin, Zhirong Yu, and Simon A. Lee. A case study exploring the current landscape of synthetic medical record generation with commercial llms. pages 105–129.
- Chang Liu, Jieshi Chen, Lee H Harrison, and Artur Dubrawski. Bridging the utility gap between maldi-tof and wgs for affordable outbreak cluster detection. pages 557–572.
- Leopoldo Julian Lechuga Lopez, Shaza Elsharief, Dhiyaa Al Jorf, Firas Darwish, Congbo Ma, and Farah E. Shamout. Uncertainty quantification for machine learning in healthcare: A survey. pages 862–907.
- Mahdi Moghaddami, Clayton Schubring, and Mohammad Siadat. Transformer model for alzheimer’s disease progression prediction using longitudinal visit sequences. pages 804–816.
- Sumit Mukherjee, ZACHARY R MCCAW, David Amar, Rounak Dey, Thomas W Soare, Hari Somineni, Nicholas Eriksson, and Colm O’Dushlaine. Causal considerations can determine the utility of machine learning assisted gwas. pages 179–193.
- Shahriar Noroozizadeh, Pim Welle, Jeremy Weiss, and George H. Chen. The impact of medication non-adherence on adverse outcomes: Evidence from schizophrenia patients via survival analysis. pages 573–609.
- YongKyung Oh and Alex Bui. Multi-view contrastive learning for robust domain adaptation in medical time series analysis. pages 502–526.

- OpenReview. Openreview, 2025. URL <https://openreview.net/>. Accessed: 2025-06-18.
- Jaeho Park, TaeJun Park, Joon-myung Kwon, and Yong-Yeon Jo. Benchmarking ecg delineation using deep neural network-based semantic segmentation models. pages 63–88.
- Arvind Pillai, Dimitris Spathis, Subigya Nepal, Amanda C. Collins, Daniel M Mackin, Michael V. Heinz, Tess Z Griffin, Nicholas C. Jacobson, and Andrew Campbell. Beyond prompting: Time2lang - bridging time-series foundation models and large language models for health sensing. pages 268–288.
- Willa Potosnak, Cristian Ignacio Challu, Kin G. Olivares, Keith A Dufendach, and Artur Dubrawski. Global deep forecasting with patient-specific pharmacokinetics. pages 306–336.
- Maya Stemmer, Lior Ungar, Talia Friedman, Lihi Bik, Yotam Hadari, Itamar Efrati, Yarden Rachamim, Lior Carmi, and Shai Fine. Predicting health states of patients with chronic pain from cellphone usage data. pages 443–457.
- Max I Sunog, Colin Magdamo, Marie-Laure Charpignon, and Mark W. Albers. Investigating primary care indications to improve the quality of electronic health record data in target trial emulation for dementia. pages 610–648.
- Jialu Tang, Tong Xia, Yuan Lu, Cecilia Mascolo, and Aaqib Saeed. Electrocardiogram–language model for few-shot question answering with meta learning. pages 89–104.
- Adedolapo Aishat Toye, Asuman Celik, and Samantha Kleinberg. Benchmarking missing data imputation methods for time series using real-world test cases. pages 480–501.
- Yoanna Turura, Sam Freesun Friedman, Aurora Cremer, Mahnaz Maddah, and Sana Tonekaboni. The latentverse: An open-source benchmarking toolkit for evaluating latent representations. pages 708–719.
- Tsai-Ning Wang, Lin-Lin Chen, Neil Zeghidour, and Aaqib Saeed. Careaqa: A cardiac and respiratory audio question answering model for open-ended diagnostic reasoning. pages 231–246, a.
- Will Ke Wang, Bill Chen, Jiamu Yang, Hayoung Jeong, Leeor Hershkovich, Shekh Md Mahmudul Islam, Mengde Liu, Ali R Roghanizad, Md Mobashir Hasan Shandhi, Andrew R Spector, and Jessilyn Dunn. Watchsleepnet: A novel model and pretraining approach for advancing sleep staging with smartwatches. pages 145–165, b.
- Shashank Yadav and Vignesh Subbian. When attention fails: Pitfalls of attention-based model interpretability for high-dimensional clinical time-series. pages 289–305.
- Mengying Yan, Meng Xia, Wei Angel Huang, Chuan Hong, Benjamin Goldstein, and Matthew M. Engelhard. Predicting partially observed long-term outcomes with adversarial positive-unlabeled domain adaptation. pages 672–690.
- Zeyu Yang, Han Yu, and Akane Sano. Contrastive pretraining for stress detection with multimodal wearable sensor data and surveys. pages 166–178.
- Leisheng Yu, Yanxiao Cai, Minxing Zhang, and Xia Hu. Self-explaining hypergraph neural networks for diagnosis prediction. pages 908–924.
- Hye Sun Yun, Karen Y.C. Zhang, Ramez Kouzy, Iain James Marshall, Junyi Jessie Li, and Byron C Wallace. Caught in the web of words: Do llms fall for spin in medical literature? pages 458–479.
- Xiao Yu Cindy Zhang, Carlos R. Ferreira, Francis Rossignol, Raymond T. Ng, Wyeth Wasserman, and Jian Zhu. Casereportbench: An llm benchmark dataset for dense information extraction in clinical case reports. pages 527–542, a.
- Xiaoman Zhang, Julian Nicolas Acosta, Hong-Yu Zhou, and Pranav Rajpurkar. Uncovering knowledge gaps in radiology report generation models through knowledge graphs. pages 30–42, b.